



Twitter Sentiment Analysis

Group 13: Chionh Wan Sim, Hiew Shani, Kushioka Nodoka, Yeo Fu Kai Marcus

OUR TEAM



Marcus

Y2 Data Science & Analytics



Shani

Y2 Data Science & Economics



Wan Sim

Y2 Data Science & Analytics



Nodoka

Y2 Data Science & Analytics

Flow of Presentation

1

Nodoka

1. Introduction of the topic and dataset
2. Issues with dataset

2

Marcus

1. Data Preprocessing
2. Random Forest
3. Logistic Regression

3

Shani

1. k-Nearest Neighbours
2. Hyperparameter Tuning
3. Evaluation Metrics

4

Wan Sim

1. Results
2. Future Improvements

Why?

Understand public sentiment on diverse subjects



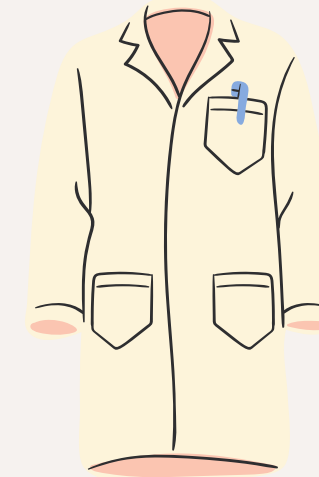
1. Business Owners

Gauge customer satisfaction and receive feedback



2. Politicians

Gauge sentiment on policies and understand societal concerns



3. Researchers

Especially useful in public health research

Reports Used and their Limitations

1. “Analytics of machine learning-based algorithms for text classification,”

✓ Clear comparison on algorithms

✗ Little details on preprocessing

2. “A Comparative Analysis of Machine Learning Classifiers for Twitter Sentiment Analysis,”

✓ Detailed pre-processing steps

✗ Trained on minimal data points

3. “Sentiment Analysis of Twitter Data,”

✓ Multiple approaches, including non-machine learning

✗ Why is the prediction accuracy of models is highly dependent on the classifiers that were taught by the target domain?

An Introduction

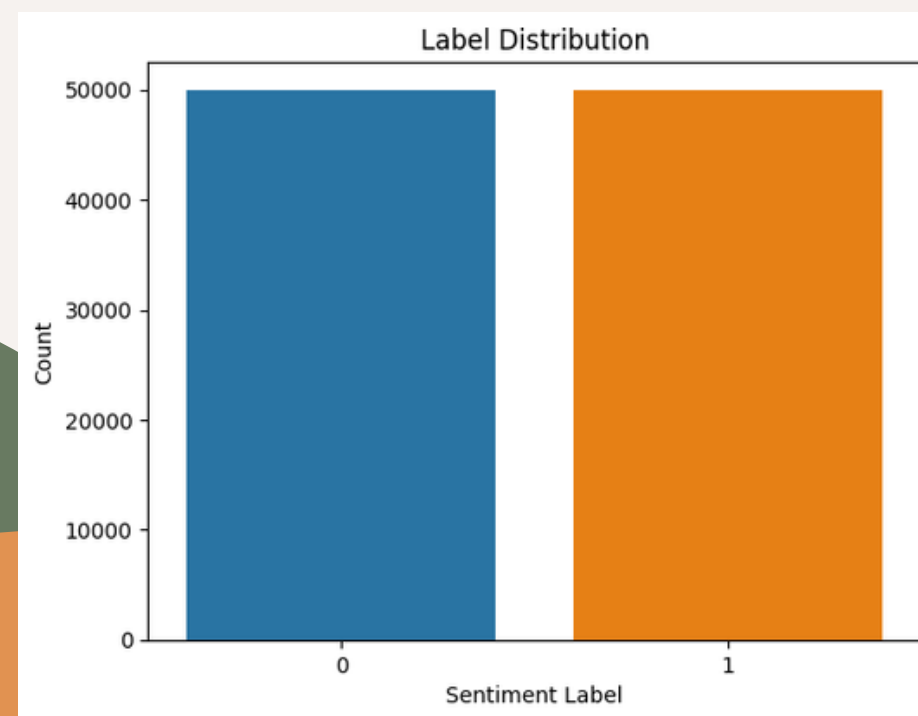
100,000 tweets



Sentiment label

1: 'Positive sentiment'

0: 'Negative Sentiment'

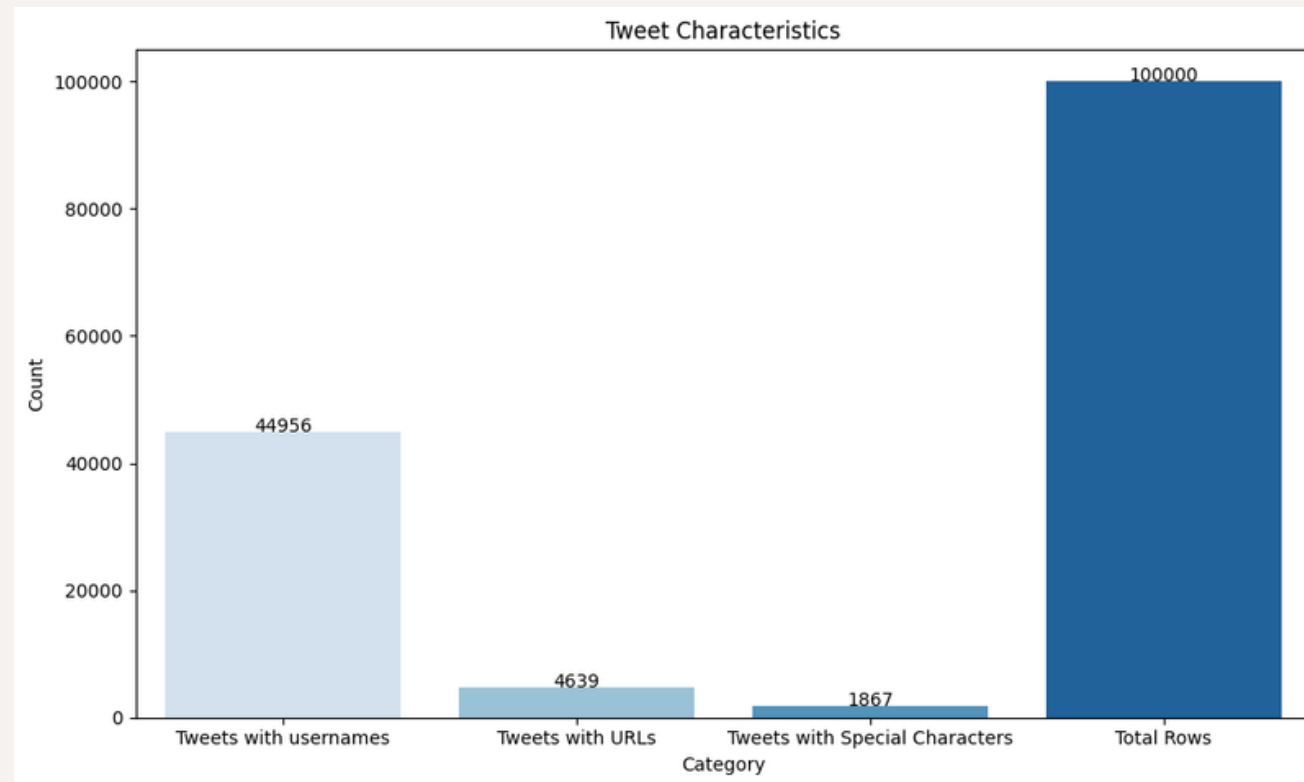


Equal distribution

50,000 instances for
both positive and
negative numbers

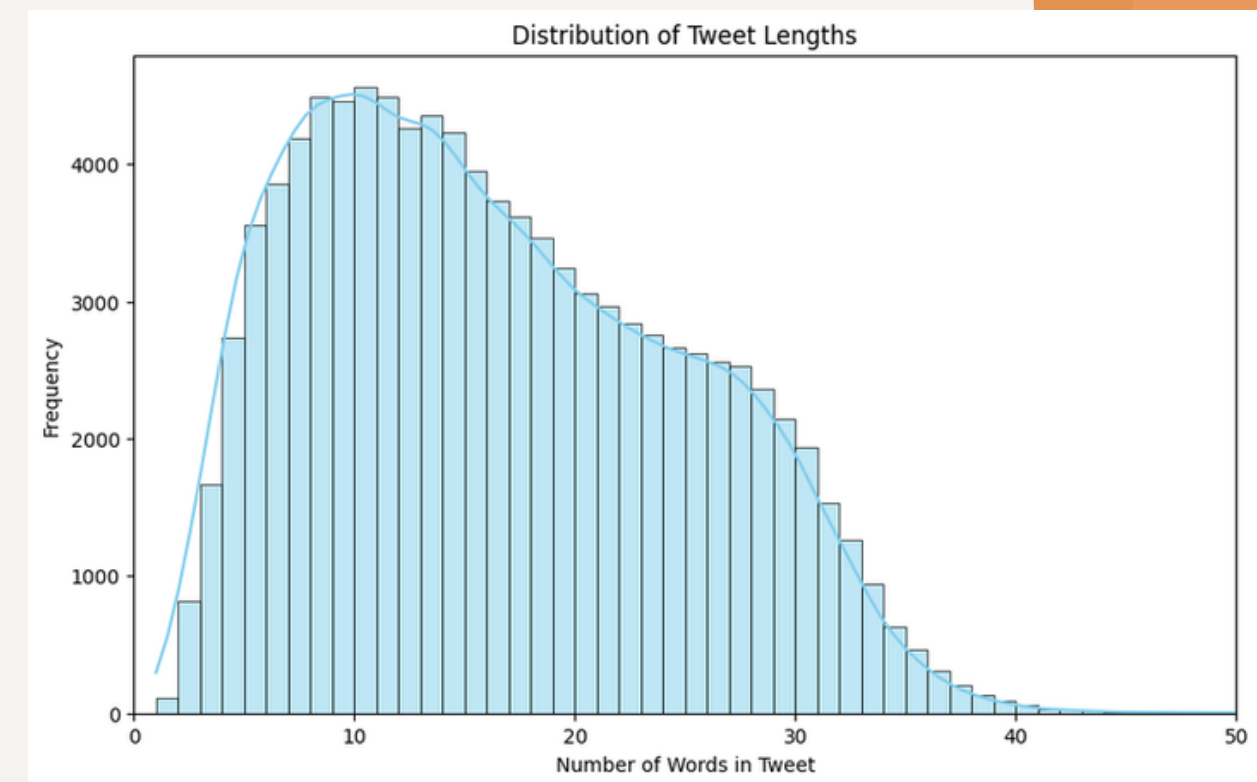


Limitations



1. Noise

Many tweets with irrelevant text



2. Short tweets

Too short to understand the context

3. Slangs and Abbreviations

NLP struggle with shorthand text

Data Preprocessing

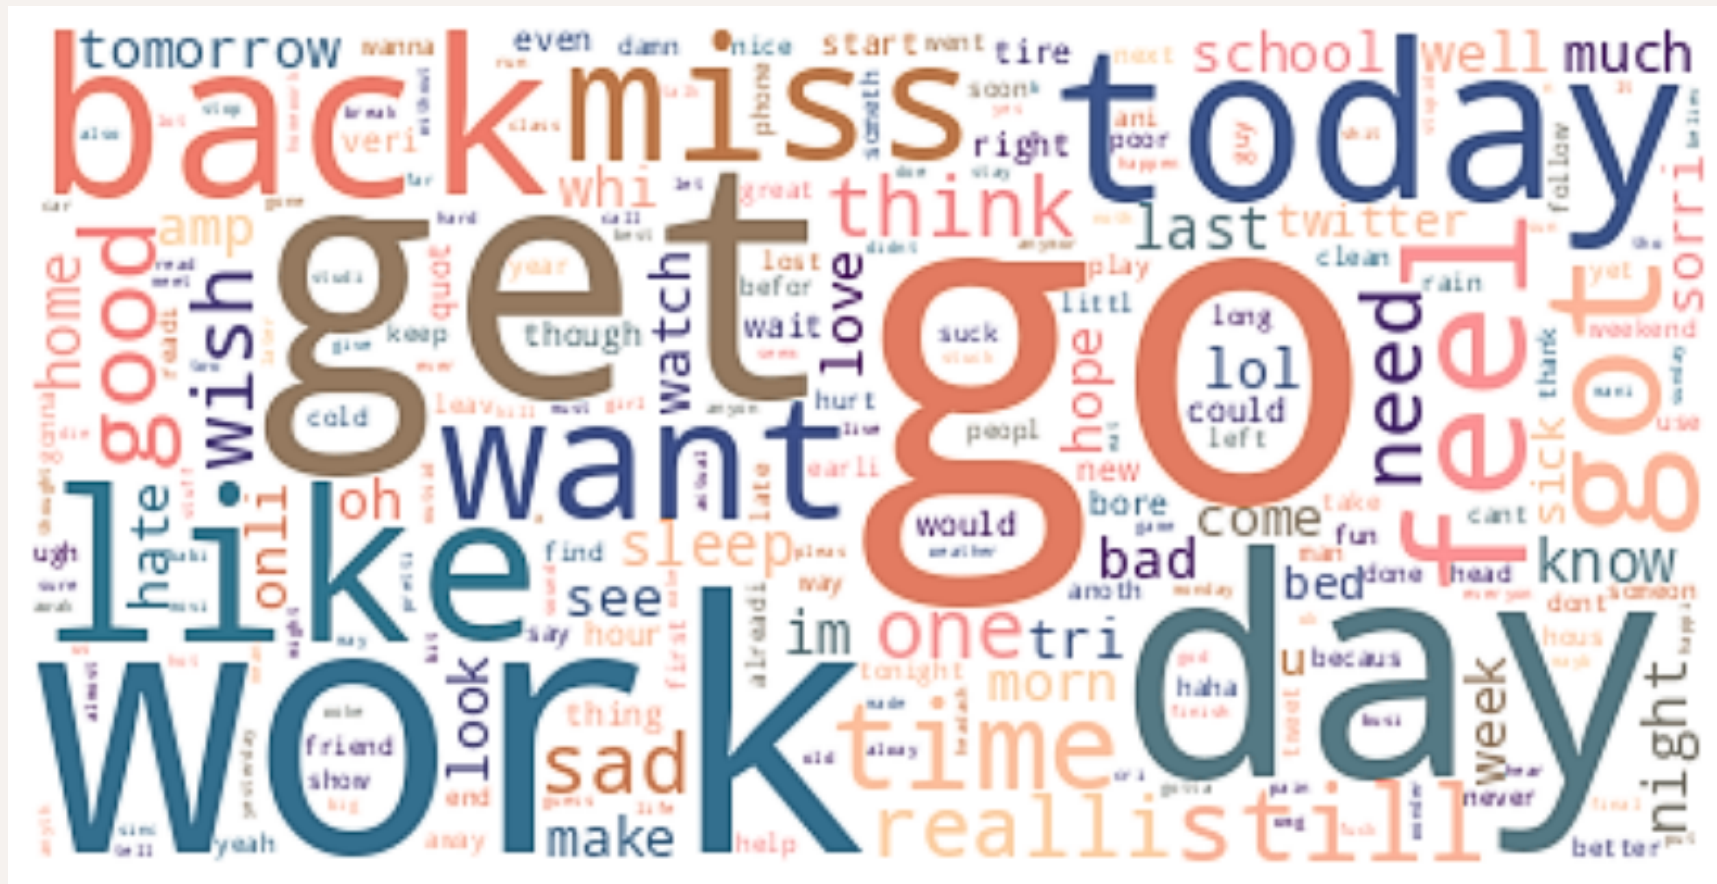
No.	Steps
1	Removed Usernames
2	Removed URLs
3	Stripped Punctuations, Numbers, and Special Characters
4	Tokenization
5	Stemming and Lemmatization
6	(Remove stop words)

“@switchfoot <http://twitpic.com/2y1zl> - Awww, that's a bummer. You shoulda got David Carr of Third Day to do it. ;D”



[david, day, got, carr, bummer, shoulda, third]

Word Cloud After Preprocessing



“Negative” Sentiments



“Positive” Sentiments

`sklearn.model_selection.train_test_split`

```
sklearn.model_selection.train_test_split(*arrays, test_size=None, train_size=None, random_state=None,  
shuffle=True, stratify=None)
```

[\[source\]](#)

Splitting into
Training &
Test Set

80%

Training Set

20%

Test Set

Tf-idf Vectorizer

`sklearn.feature_extraction.text.TfidfVectorizer`

```
class sklearn.feature_extraction.text.TfidfVectorizer(*, input='content', encoding='utf-8',  
decode_error='strict', strip_accents=None, lowercase=True, preprocessor=None, tokenizer=None, analyzer='word',  
stop_words=None, token_pattern='(?u)|b|w|w+|b', ngram_range=(1, 1), max_df=1.0, min_df=1, max_features=None,  
vocabulary=None, binary=False, dtype=<class 'numpy.float64'>, norm='l2', use_idf=True, smooth_idf=True,  
sublinear_tf=False)
```

[\[source\]](#)

```
[ ] vectorizer = TfidfVectorizer(ngram_range=(1, 3), max_features = 10000, stop_words='english')
```

Takes into account
unigrams, bigrams, and trigrams

Analyse the top 10,000
most frequent terms

Remove common
English stop words

Fit Transform

```
fit_transform(raw_documents, y=None)
```

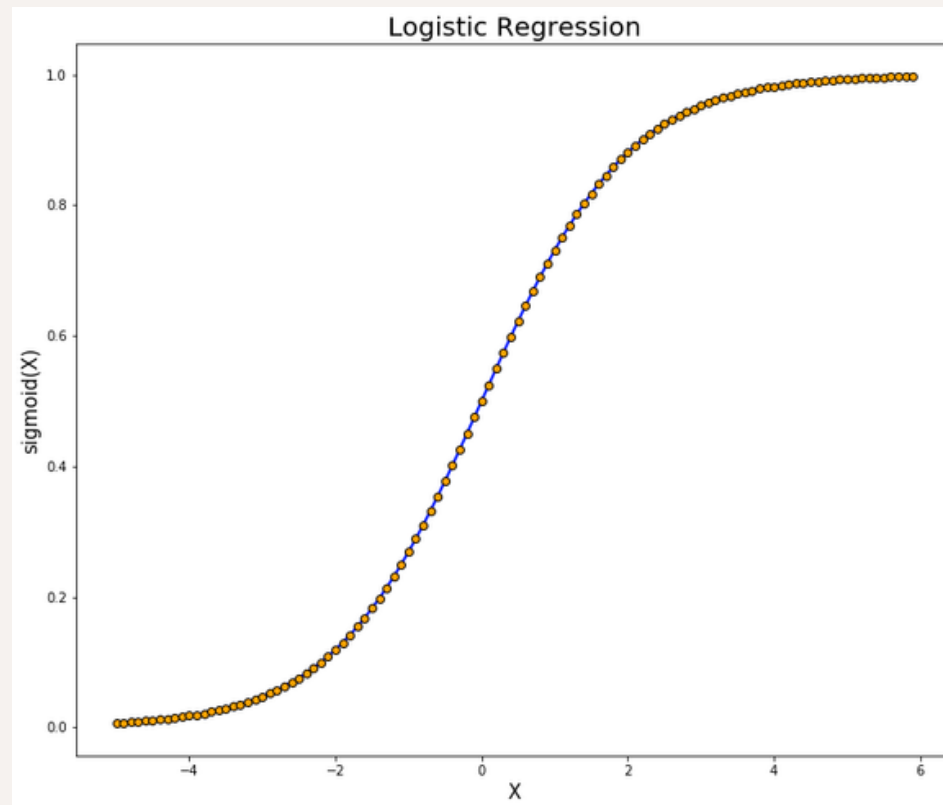
[\[source\]](#)

Learn the vocabulary dictionary and return document-term matrix.

This is equivalent to fit followed by transform, but more efficiently implemented.

Parameters:	raw_documents : <i>iterable</i> An iterable which generates either str, unicode or file objects.
	y : <i>None</i> This parameter is ignored.
Returns:	X : <i>array of shape (n_samples, n_features)</i> Document-term matrix.

Logistic Regression



`sklearn.linear_model.LogisticRegression`

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

[\[source\]](#)

What is Logistics Regression?

A statistical model for binary classification, estimates the probability of a data point belonging to a specific class through a logistic function



Simplistic



Computational
Efficient



Reveals most
influential words

Key Parameters

C : float, default=1.0

Inverse of regularization strength; must be a positive float. Like in support vector machines, smaller values specify stronger regularization.

Random Forest

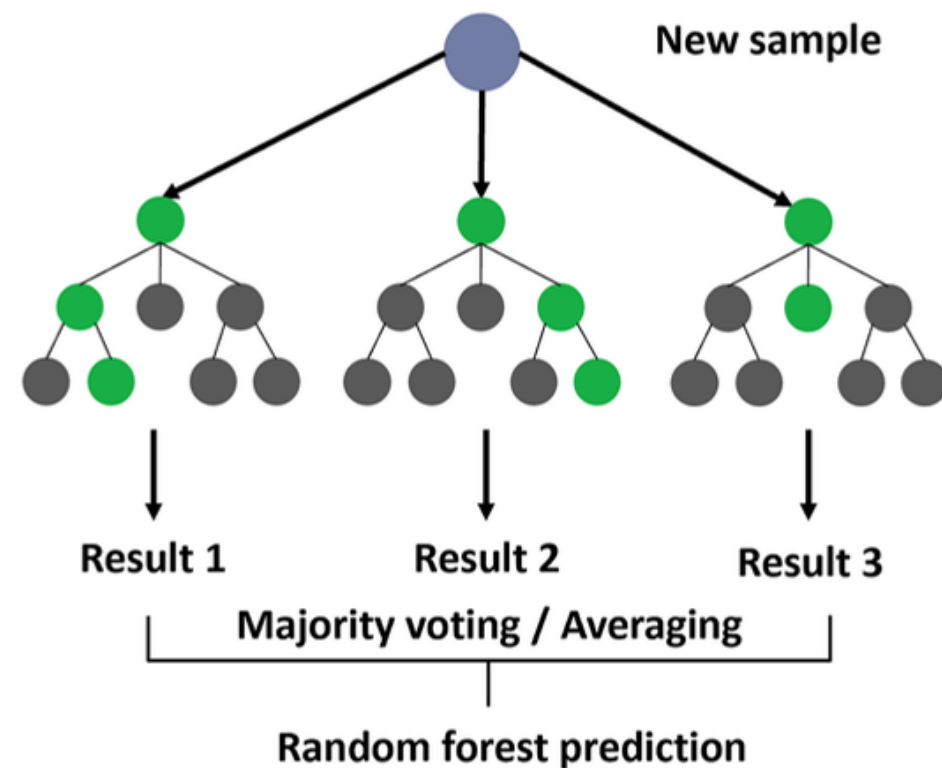
`sklearn.ensemble.RandomForestClassifier`

```
class sklearn.ensemble.RandomForestClassifier(n_estimators=100, *, criterion='gini', max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='sqrt', max_leaf_nodes=None, min_impurity_decrease=0.0, bootstrap=True, oob_score=False, n_jobs=None, random_state=None, verbose=0, warm_start=False, class_weight=None, ccp_alpha=0.0, max_samples=None)
```

[source]

What is Random Forest?

It is an ensemble learning method. It constructs multiple decision trees and outputs the mode class



Ability to handle large, high-dimensional datasets



Robustness against overfitting

Key Parameters

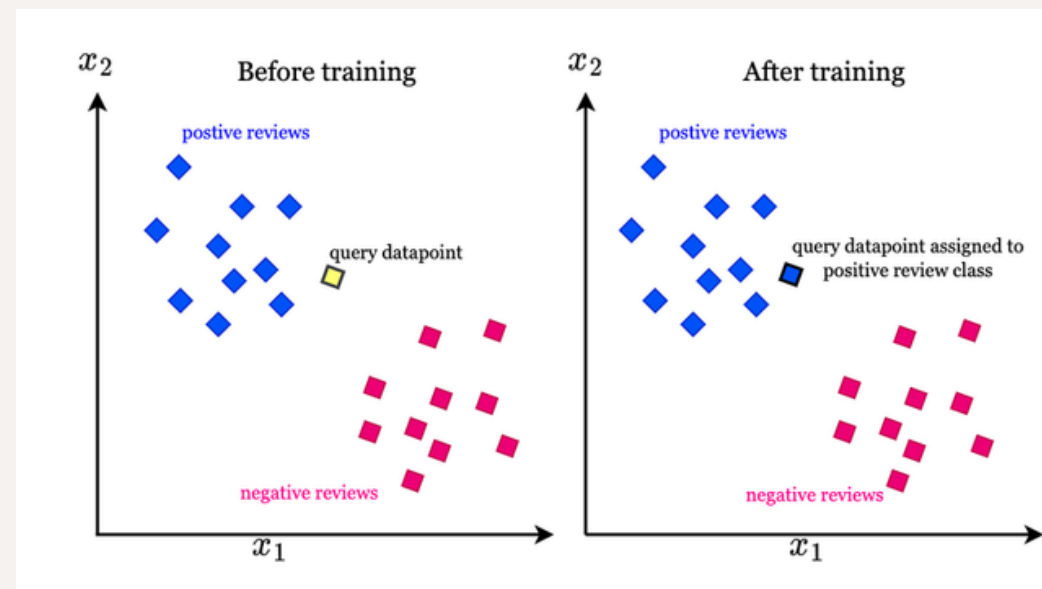
`max_depth` : *int*, *default=None*

The maximum depth of the tree. If `None`, then nodes are expanded until all leaves are pure or until all leaves contain less than `min_samples_split` samples.

`n_estimators` : *int*, *default=100*

The number of trees in the forest.

K Nearest Neighbours (KNN)



`sklearn.neighbors.KNeighborsClassifier` ¶

```
class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=None)
```

[\[source\]](#)

What is K-Nearest Neighbour?

An instance-based learning algorithm that classifies data points based on the majority class among their k nearest neighbours in a feature space



Makes no explicit assumptions



Flexibility in modelling complex relationships

Key Parameters

`n_neighbors` : *int*, *default=5*

Number of neighbors to use by default for `kneighbors` queries.

`p` : *float*, *default=2*

Power parameter for the Minkowski metric. When $p = 1$, this is equivalent to using `manhattan_distance` (l_1), and `euclidean_distance` (l_2) for $p = 2$. For arbitrary p , `minkowski_distance` (l_p) is used.

Hyperparameter Tuning



RandomizedSearchCV

`sklearn.grid_search.RandomizedSearchCV`



```
class sklearn.grid_search.RandomizedSearchCV(estimator, param_distributions, n_iter=10,  
scoring=None, fit_params=None, n_jobs=1, iid=True, refit=True, cv=None, verbose=0,  
pre_dispatch='2*n_jobs', random_state=None, error_score='raise')
```

[\[source\]](#)



Faster: only covers a subset of hyperparameter combinations



Risk of Missing Optimal Values



GridSearchCV

`sklearn.model_selection.GridSearchCV`

```
class sklearn.model_selection.GridSearchCV(estimator, param_grid, *,  
scoring=None, n_jobs=None, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs',  
error_score=nan, return_train_score=False)
```

[\[source\]](#)



Exhaustively explores all parameter combinations



Computationally expensive

Hyperparameter Tuning

RandomizedSearchCV

`sklearn.grid_search.RandomizedSearchCV`

```
class sklearn.grid_search.RandomizedSearchCV(estimator, param_distributions, n_iter=10,  
scoring=None, fit_params=None, n_jobs=1, iid=True, refit=True, cv=None, verbose=0,  
pre_dispatch='2*n_jobs', random_state=None, error_score='raise')
```

[\[source\]](#)



Faster: only covers a subset of hyperparameter combinations



Risk of Missing Optimal Values

GridSearchCV

`sklearn.model_selection.GridSearchCV`

```
class sklearn.model_selection.GridSearchCV(estimator, param_grid, *,  
scoring=None, n_jobs=None, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs',  
error_score=nan, return_train_score=False)
```

[\[source\]](#)

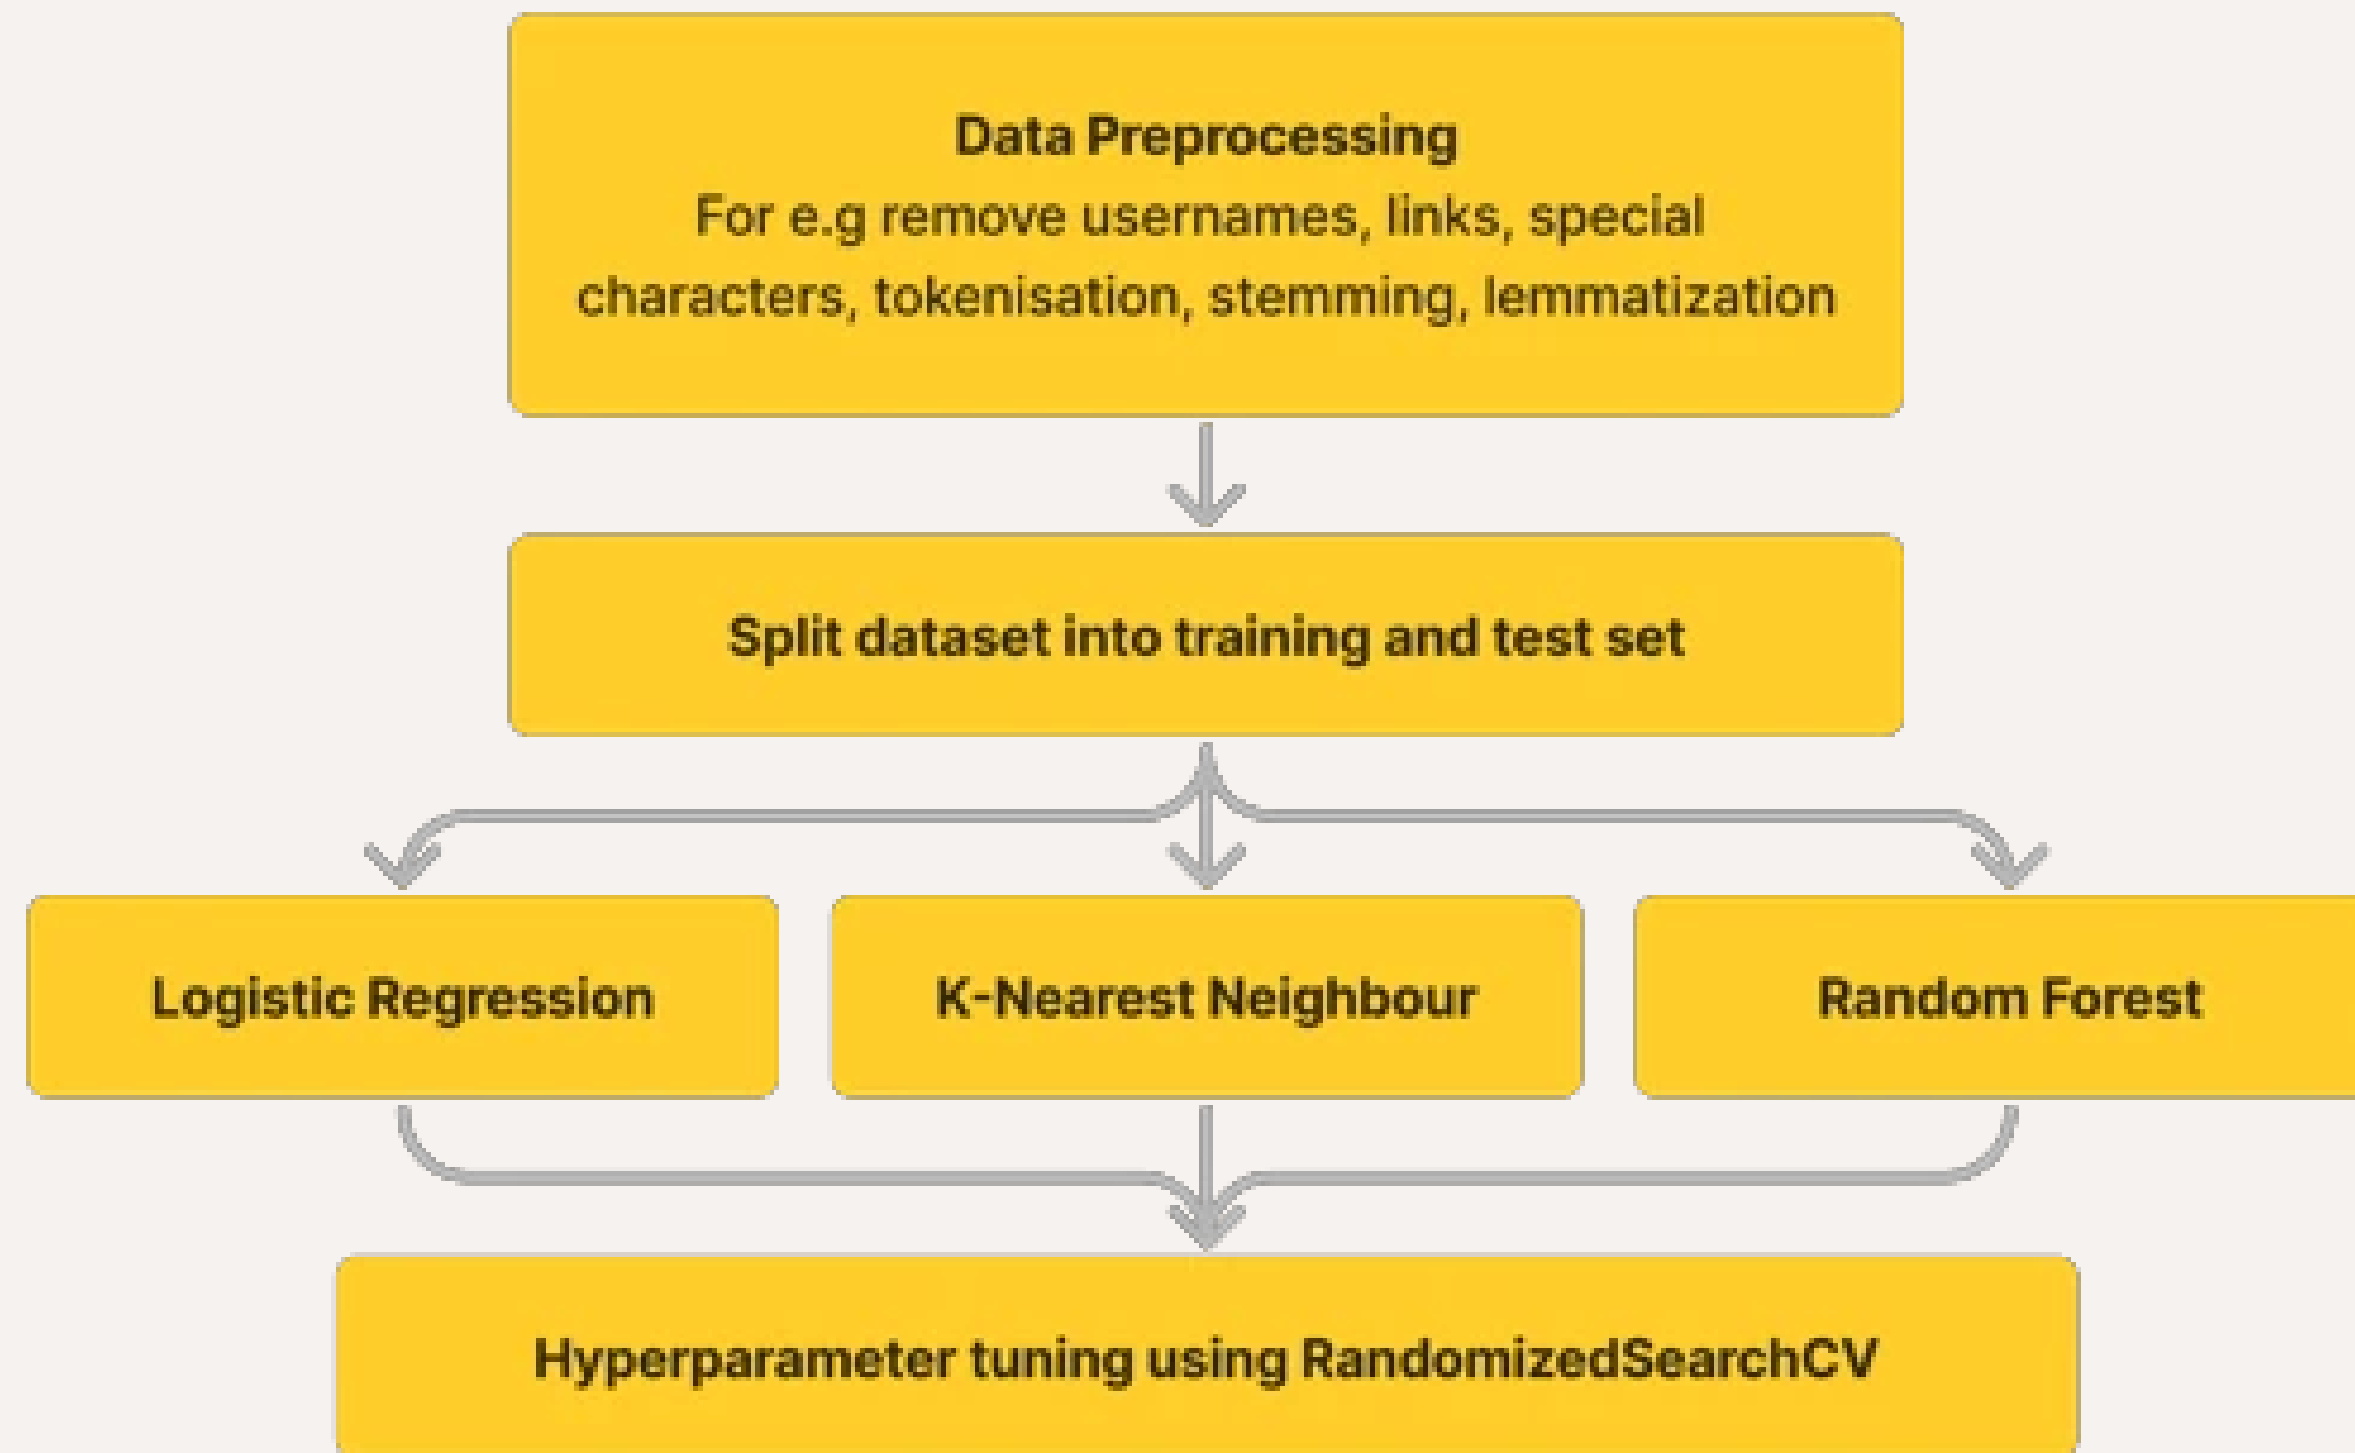


Exhaustively explores all parameter combinations



Computationally expensive

Overall Methodology



Evaluation Metrics

Accuracy

$$\frac{TP+TN}{TP+TN+FP+FN}$$

a measure of the overall correctness of a classification model

Precision

$$\frac{TP}{TP+FP}$$

assesses the model's accuracy when it predicts a positive class

Recall

$$\frac{TP}{TP+FN}$$

measures the model's ability to correctly identify all relevant instances of a specific class.

F1-Score

$$\frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

It provides a balance between precision and recall, rewarding models that have both high precision and high recall.

Evaluation Metrics

Accuracy

$$\frac{TP+TN}{TP+TN+FP+FN}$$

a measure of the overall correctness of a classification model

Precision

$$\frac{TP}{TP+FP}$$

assesses the model's accuracy when it predicts a positive class

Recall

$$\frac{TP}{TP+FN}$$

measures the model's ability to correctly identify all relevant instances of a specific class.

F1-Score

$$\frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

It provides a balance between precision and recall, rewarding models that have both high precision and high recall.

Results

Model: Logistic Regression

Best Hyperparameters: {'C': 1}

Accuracy on Validation Set: 0.75



Evaluation on Test Set:

Accuracy on Test Set: 0.75

Classification Report on Test Set:

	precision	recall	f1-score	support
Negative	0.77	0.73	0.75	10037
Positive	0.74	0.78	0.76	9963
accuracy			0.75	20000
macro avg	0.75	0.75	0.75	20000
weighted avg	0.75	0.75	0.75	20000

Results

Model: Random Forest

Best Hyperparameters: {'n_estimators': 175, 'max_depth': None}

Accuracy on Validation Set: 0.74

F1 Score on Validation Set: 0.74

Evaluation on Test Set:

Accuracy on Test Set: 0.74

F1 Score on Test Set: 0.74

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.73	0.74	10037
1	0.74	0.75	0.74	9963
accuracy			0.74	20000
macro avg	0.74	0.74	0.74	20000
weighted avg	0.74	0.74	0.74	20000

Results

Model: K-Nearest Neighbors

Best Hyperparameters: {'n_neighbors': 2}

Accuracy on Validation Set: 0.61



Evaluation on Test Set:

Accuracy on Test Set: 0.61

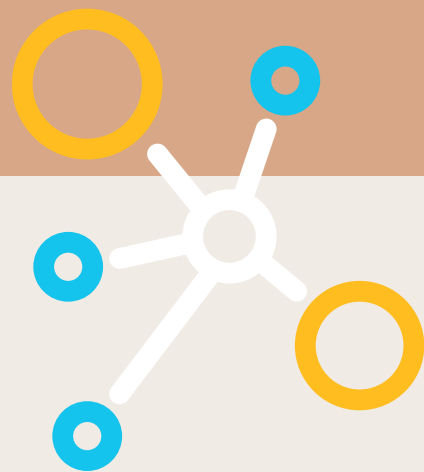
Classification Report on Test Set:

	precision	recall	f1-score	support
Negative	0.57	0.85	0.68	10037
Positive	0.71	0.36	0.48	9963
accuracy			0.61	20000
macro avg	0.64	0.61	0.58	20000
weighted avg	0.64	0.61	0.58	20000

Future Improvements

Alternative Models

Ensemble Methods
Deep Learning models:
Recurrent Neural
Networks (RNNs)



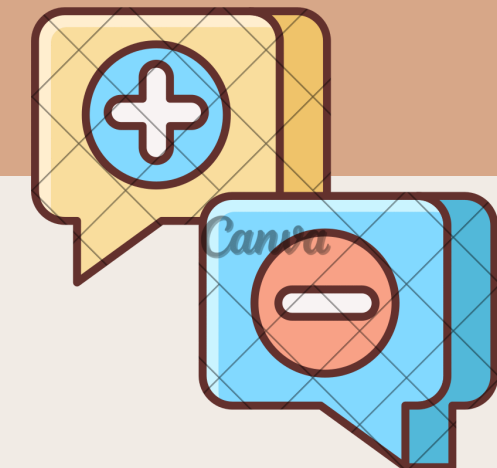
Data Manipulation

Oversampling
Undersampling
Synthetic data generation



Additional Information

incorporating annotated
lists of words with their
sentiment polarity



Q & A



Workload Breakdown



Marcus

40%



Shani

30%



Wan Sim

10%



Nodoka

20%